

S.Kavin Siva

3122215001306

CSE-A

UCS2612 Machine Learning Laboratory

ASSIGNMENT 8 : Applications of Random Forest and AdaBoost Ensemble Techniques

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

LOADING DATASET

```
url =
"https://archive.ics.uci.edu/ml/machine-learning-databases/breast-
cancer-wisconsin/wdbc.data"
names = ['ID', 'Diagnosis', 'mean_radius', 'mean_texture', 'mean_perimeter',
'mean_area',
        'mean_smoothness', 'mean_compactness', 'mean_concavity',
'mean_concave_points',
        'mean_symmetry', 'mean_fractal_dimension', 'se_radius', 'se_texture',
        'se_perimeter', 'se_area', 'se_smoothness', 'se_compactness',
'se_concavity',
        'se_concave_points', 'se_symmetry', 'se_fractal_dimension',
'worst_radius',
        'worst_texture', 'worst_perimeter', 'worst_area', 'worst_smoothness',
        'worst_compactness', 'worst_concavity',
'worst_concave_points',
        'worst_symmetry', 'worst_fractal_dimension']data
= pd.read_csv(url, names=names)
```

data

	ID	Diagnosis	mean_radius	mean_texture	mean_perimeter
mean_area \					
0	842302	M	17.99	10.38	122.80
1001.0					
1	842517	M	20.57	17.77	132.90
1326.0					
2	84300903	M	19.69	21.25	130.00
1203.0					
3	84348301	M	11.42	20.38	77.58
386.1					
4	84358402	M	20.29	14.34	135.10
1297.0					

564	926424	M	21.56	22.39	142.00
1479.0					
565	926682	M	20.13	28.25	131.20
1261.0					
566	926954	M	16.60	28.08	108.30
858.1					
567	927241	M	20.60	29.33	140.10
1265.0					
568	92751	B	7.76	24.54	47.92
181.0					

	mean_smoothness	ean_compactness	mean_concavity
mean_concave_points \			
0	0.11840	0.27760	0.30010
0.14710			
1	0.08474	0.07864	0.08690
0.07017			
2	0.10960	0.15990	0.19740
0.12790			
3	0.14250	0.28390	0.24140
0.10520			
4	0.10030	0.13280	0.19800
0.10430			
..
...			
564	0.11100	0.11590	0.24390
0.13890			
565	0.09780	0.10340	0.14400
0.09791			
566	0.08455	0.10230	0.09251
0.05302			
567	0.11780	0.27700	0.35140
0.15200			
568	0.05263	0.04362	0.00000
0.00000			

	...	worst_radius	worst_texture	worst_perimeter	worst_area	\
0	...	25.380	17.33	184.60	2019.0	
1	...	24.990	23.41	158.80	1956.0	
2	...	23.570	25.53	152.50	1709.0	
3	...	14.910	26.50	98.87	567.7	
4	...	22.540	16.67	152.20	1575.0	
..	
564	...	25.450	26.40	166.10	2027.0	
565	...	23.690	38.25	155.00	1731.0	
566	...	18.980	34.12	126.70	1124.0	
567	...	25.740	39.42	184.60	1821.0	
568	...	9.456	30.37	59.16	268.6	

worst_smoothness worst_compactness worst_concavity \

0	0.16220	0.66560	0.7119
1	0.12380	0.18660	0.2416
2	0.14440	0.42450	0.4504
3	0.20980	0.86630	0.6869
4	0.13740	0.20500	0.4000
..
564	0.14100	0.21130	0.4107
565	0.11660	0.19220	0.3215
566	0.11390	0.30940	0.3403
567	0.16500	0.86810	0.9387
568	0.08996	0.06444	0.0000

	worst_concave_points	worst_symmetry	worst_fractal_dimension
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678
..
564	0.2216	0.2060	0.07115
565	0.1628	0.2572	0.06637
566	0.1418	0.2218	0.07820
567	0.2650	0.4087	0.12400
568	0.0000	0.2871	0.07039

[569 rows x 32 columns]

data.head()

	ID	Diagnosis	mean_radius	mean_texture	mean_perimeter
mean area \					
0	842302	M	17.99	10.38	122.80
1001.0					
1	842517	M	20.57	17.77	132.90
1326.0					
2	84300903	M	19.69	21.25	130.00
1203.0					
3	84348301	M	11.42	20.38	77.58
386.1					
4	84358402	M	20.29	14.34	135.10
1297.0					

	mean_smoothness	mean_compactness	mean_concavity
mean concave points \			
0	0.11840	0.27760	0.3001
0.14710			
1	0.08474	0.07864	0.0869
0.07017			
2	0.10960	0.15990	0.1974
0.12790			

```

3          0.14250          0.28390          0.2414
0.10520
4          0.10030          0.13280          0.1980
0.10430

```

```

...  worst_radius  worst_texture  worst_perimeter  worst_area  \
0  ...      25.38      17.33      184.60      2019.0
1  ...      24.99      23.41      158.80      1956.0
2  ...      23.57      25.53      152.50      1709.0
3  ...      14.91      26.50       98.87       567.7
4  ...      22.54      16.67      152.20      1575.0

```

```

      worst_smoothness  worst_compactness  worst_concavity
worst_concave_points  \
0          0.1622          0.6656          0.7119
0.2654
1          0.1238          0.1866          0.2416
0.1860
2          0.1444          0.4245          0.4504
0.2430
3          0.2098          0.8663          0.6869
0.2575
4          0.1374          0.2050          0.4000
0.1625

```

```

      worst_symmetry  worst_fractal_dimension
0          0.4601          0.11890
1          0.2750          0.08902
2          0.3613          0.08758
3          0.6638          0.17300
4          0.2364          0.07678

```

```
[5 rows x 32 columns]
```

```
data.describe()
```

```

              ID mean_radius mean_texture mean_perimeter
mean_area  \
count  5.690000e+02   569.000000   569.000000   569.000000
569.000000
mean    3.037183e+07   14.127292   19.289649   91.969033
654.889104
std     1.250206e+08    3.524049    4.301036   24.298981
351.914129
min     8.670000e+03    6.981000    9.710000   43.790000
143.500000
25%     8.692180e+05   11.700000   16.170000   75.170000
420.300000
50%     9.060240e+05   13.370000   18.840000   86.240000
551.100000

```

75%	8.813129e+06	15.780000	21.800000	104.100000
782.700000				
max	9.113205e+08	28.110000	39.280000	188.500000
2501.000000				

	mean_smoothness	mean_compactness	mean_concavity
mean_concave_points \			
count	569.000000	569.000000	569.000000
569.000000			
mean	0.096360	0.104341	0.088799
0.048919			
std	0.014064	0.052813	0.079720
0.038803			
min	0.052630	0.019380	0.000000
0.000000			
25%	0.086370	0.064920	0.029560
0.020310			
50%	0.095870	0.092630	0.061540
0.033500			
75%	0.105300	0.130400	0.130700
0.074000			
max	0.163400	0.345400	0.426800
0.201200			

	mean_symmetry	...	worst_radius	worst_texture
worst_perimeter \				
count	569.000000	...	569.000000	569.000000
569.000000				
mean	0.181162	...	16.269190	25.677223
107.261213				
std	0.027414	...	4.833242	6.146258
33.602542				
min	0.106000	...	7.930000	12.020000
50.410000				
25%	0.161900	...	13.010000	21.080000
84.110000				
50%	0.179200	...	14.970000	25.410000
97.660000				
75%	0.195700	...	18.790000	29.720000
125.400000				
max	0.304000	...	36.040000	49.540000
251.200000				

	worst_area	worst_smoothness	worst_compactness
worst_concavity \			
count	569.000000	569.000000	569.000000
569.000000			
mean	880.583128	0.132369	0.254265
0.272188			
std	569.356993	0.022832	0.157336

```

0.208624
min      185.200000      0.071170      0.027290
0.000000
25%      515.300000      0.116600      0.147200
0.114500
50%      686.500000      0.131300      0.211900
0.226700
75%     1084.000000      0.146000      0.339100
0.382900
max     4254.000000      0.222600      1.058000
1.252000

```

	worst_concave_points	worst_symmetry	worst_fractal_dimension
count	569.000000	569.000000	569.000000
mean	0.114606	0.290076	0.083946
std	0.065732	0.061867	0.018061
min	0.000000	0.156500	0.055040
25%	0.064930	0.250400	0.071460
50%	0.099930	0.282200	0.080040
75%	0.161400	0.317900	0.092080
max	0.291000	0.663800	0.207500

```
[8 rows x 31 columns]
```

```
num_rows, num_columns = data.shape
```

```
print("Number of rows:", num_rows)
print("Number of columns:", num_columns)
```

```
Number of rows: 569
Number of columns: 32
```

```
data.nunique()
```

ID	569
Diagnosis	2
mean_radius	456
mean_texture	479
mean_perimeter	522
mean_area	539
mean_smoothness	474
mean_compactness	537
mean_concavity	537
mean_concave_points	542
mean_symmetry	432
mean_fractal_dimension	499
se_radius	540
se_texture	519
se_perimeter	533
se_area	528

se_smoothness	547
se_compactness	541
se_concavity	533
se_concave_points	507
se_symmetry	498
se_fractal_dimension	545
worst_radius	457
worst_texture	511
worst_perimeter	514
worst_area	544
worst_smoothness	411
worst_compactness	529
worst_concavity	539
worst_concave_points	492
worst_symmetry	500
worst_fractal_dimension	535
dtype:	int64

PRE-PROCESSING DATA

1.HANDLING MISSING VALUES

```
print("The Number of Missing Values in the dataset\n")  
data.isnull().sum()
```

```
The Number of Missing Values in the dataset
```

```
worst_radius      0
worst_texture     0
worst_perimeter   0
worst_area        0
worst_smoothness  0
worst_compactness 0
worst_concavity   0
worst_concave_points 0
worst_symmetry    0
worst_fractal_dimension 0
dtype: int64
```

2. ENCODING CATEGORICAL TARGET VARIABLE

```
label_encoder = LabelEncoder()
data['Diagnosis'] = label_encoder.fit_transform(data['Diagnosis'])
```

data

	ID	Diagnosis	mean_radius	mean_texture	mean_perimeter	\
0	842302	1	17.99	10.38	122.80	
1	842517	1	20.57	17.77	132.90	
2	84300903	1	19.69	21.25	130.00	
3	84348301	1	11.42	20.38	77.58	
4	84358402	1	20.29	14.34	135.10	
..	
564	926424	1	21.56	22.39	142.00	
565	926682	1	20.13	28.25	131.20	
566	926954	1	16.60	28.08	108.30	
567	927241	1	20.60	29.33	140.10	
568	92751	0	7.76	24.54	47.92	

	mean_area	mean_smoothness	mean_compactness	mean_concavity	\
0	1001.0	0.11840	0.27760	0.30010	
1	1326.0	0.08474	0.07864	0.08690	
2	1203.0	0.10960	0.15990	0.19740	
3	386.1	0.14250	0.28390	0.24140	
4	1297.0	0.10030	0.13280	0.19800	
..	
564	1479.0	0.11100	0.11590	0.24390	
565	1261.0	0.09780	0.10340	0.14400	
566	858.1	0.08455	0.10230	0.09251	
567	1265.0	0.11780	0.27700	0.35140	
568	181.0	0.05263	0.04362	0.00000	

	mean_concave_points	...	worst_radius	worst_texture
worst_perimeter \				
0	0.14710	...	25.380	17.33
184.60				
1	0.07017	...	24.990	23.41

158.80				
2	0.12790	...	23.570	25.53
152.50				
3	0.10520	...	14.910	26.50
98.87				
4	0.10430	...	22.540	16.67
152.20				
..
...				
564	0.13890	...	25.450	26.40
166.10				
565	0.09791	...	23.690	38.25
155.00				
566	0.05302	...	18.980	34.12
126.70				
567	0.15200	...	25.740	39.42
184.60				
568	0.00000	...	9.456	30.37
59.16				

	worst_area	worst_smoothness	worst_compactness	worst_concavity
\				
0	2019.0	0.16220	0.66560	0.7119
1	1956.0	0.12380	0.18660	0.2416
2	1709.0	0.14440	0.42450	0.4504
3	567.7	0.20980	0.86630	0.6869
4	1575.0	0.13740	0.20500	0.4000
..
564	2027.0	0.14100	0.21130	0.4107
565	1731.0	0.11660	0.19220	0.3215
566	1124.0	0.11390	0.30940	0.3403
567	1821.0	0.16500	0.86810	0.9387
568	268.6	0.08996	0.06444	0.0000

	worst_concave_points	worst_symmetry	worst_fractal_dimension
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300

```

4          0.1625          0.2364          0.07678
..          ...          ...          ...
564        0.2216        0.2060        0.07115
565        0.1628        0.2572        0.06637
566        0.1418        0.2218        0.07820
567        0.2650        0.4087        0.12400
568        0.0000        0.2871        0.07039

```

```
[569 rows x 32 columns]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	569 non-null	int64
1	Diagnosis	569 non-null	int32
2	mean_radius	569 non-null	float64
3	mean_texture	569 non-null	float64
4	mean_perimeter	569 non-null	float64
5	mean_area	569 non-null	float64
6	mean_smoothness	569 non-null	float64
7	mean_compactness	569 non-null	float64
8	mean_concavity	569 non-null	float64
9	mean_concave_points	569 non-null	float64
10	mean_symmetry	569 non-null	float64
11	mean_fractal_dimension	569 non-null	float64
12	se_radius	569 non-null	float64
13	se_texture	569 non-null	float64
14	se_perimeter	569 non-null	float64
15	se_area	569 non-null	float64
16	se_smoothness	569 non-null	float64
17	se_compactness	569 non-null	float64
18	se_concavity	569 non-null	float64
19	se_concave_points	569 non-null	float64
20	se_symmetry	569 non-null	float64
21	se_fractal_dimension	569 non-null	float64
22	worst_radius	569 non-null	float64
23	worst_texture	569 non-null	float64
24	worst_perimeter	569 non-null	float64
25	worst_area	569 non-null	float64
26	worst_smoothness	569 non-null	float64
27	worst_compactness	569 non-null	float64
28	worst_concavity	569 non-null	float64
29	worst_concave_points	569 non-null	float64
30	worst_symmetry	569 non-null	float64
31	worst_fractal_dimension	569 non-null	float64

```
dtypes: float64(30), int32(1), int64(1)
memory usage: 140.2 KB
```

1. NORMALIZATION AND STANDARDIZATION

```
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data.drop(['ID', 'Diagnosis'],axis=1))
data_scaled = pd.DataFrame(data_scaled, columns=data.columns[2:])
```

data

	ID	Diagnosis	mean_radius	mean_texture	mean_perimeter	\
0	842302	1	17.99	10.38	122.80	
1	842517	1	20.57	17.77	132.90	
2	84300903	1	19.69	21.25	130.00	
3	84348301	1	11.42	20.38	77.58	
4	84358402	1	20.29	14.34	135.10	
...	
564	926424	1	21.56	22.39	142.00	
565	926682	1	20.13	28.25	131.20	
566	926954	1	16.60	28.08	108.30	
567	927241	1	20.60	29.33	140.10	
568	92751	0	7.76	24.54	47.92	

	mean_area	mean_smoothness	mean_compactness	mean_concavity	\
0	1001.0	0.11840	0.27760	0.30010	
1	1326.0	0.08474	0.07864	0.08690	
2	1203.0	0.10960	0.15990	0.19740	
3	386.1	0.14250	0.28390	0.24140	
4	1297.0	0.10030	0.13280	0.19800	
...	
564	1479.0	0.11100	0.11590	0.24390	
565	1261.0	0.09780	0.10340	0.14400	
566	858.1	0.08455	0.10230	0.09251	
567	1265.0	0.11780	0.27700	0.35140	
568	181.0	0.05263	0.04362	0.00000	

	mean_concave_points	...	worst_radius	worst_texture
worst_perimeter \				
0	0.14710	...	25.380	17.33
184.60				
1	0.07017	...	24.990	23.41
158.80				
2	0.12790	...	23.570	25.53
152.50				
3	0.10520	...	14.910	26.50
98.87				
4	0.10430	...	22.540	16.67
152.20				

```

..          ...      ...          ...          ...
...
564          0.13890    ...          25.450          26.40
166.10
565          0.09791    ...          23.690          38.25
155.00
566          0.05302    ...          18.980          34.12
126.70
567          0.15200    ...          25.740          39.42
184.60
568          0.00000    ...          9.456          30.37
59.16

```

	worst_area	worst_smoothness	worst_compactness	worst_concavity
\				
0	2019.0	0.16220	0.66560	0.7119
1	1956.0	0.12380	0.18660	0.2416
2	1709.0	0.14440	0.42450	0.4504
3	567.7	0.20980	0.86630	0.6869
4	1575.0	0.13740	0.20500	0.4000

```

..          ...      ...          ...          ...
564          2027.0          0.14100          0.21130          0.4107
565          1731.0          0.11660          0.19220          0.3215
566          1124.0          0.11390          0.30940          0.3403
567          1821.0          0.16500          0.86810          0.9387
568          268.6          0.08996          0.06444          0.0000

```

	worst_concave_points	worst_symmetry	worst_fractal_dimension
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678
..
564	0.2216	0.2060	0.07115
565	0.1628	0.2572	0.06637
566	0.1418	0.2218	0.07820
567	0.2650	0.4087	0.12400
568	0.0000	0.2871	0.07039


[569 rows x 32 columns]

EXPLORATORY DATA ANALYSIS

1. CORRELATION HEATMAP

```
import matplotlib.pyplot as plt
import seaborn as sns
correlation_matrix = data.corr()


plt.figure(figsize=(25, 15))
sns.heatmap(correlation_matrix, annot=True, linecolor='black',
            fmt='.2f', linewidths=.1)
plt.title('Correlation Heatmap')
plt.show()
```



1. BAR-CHART

```
data['Diagnosis'].value_counts().plot(kind='bar', edgecolor='black', col
or=['skyblue', 'salmon'])
plt.title("Diagnosis", fontsize=20)
```

```
plt.show()  
data['Diagnosis'].value_counts()
```




```
Diagnosis  
0      357  
1      212  
Name: count, dtype: int64
```

1. PIE-CHART OF TARGET COLUMN

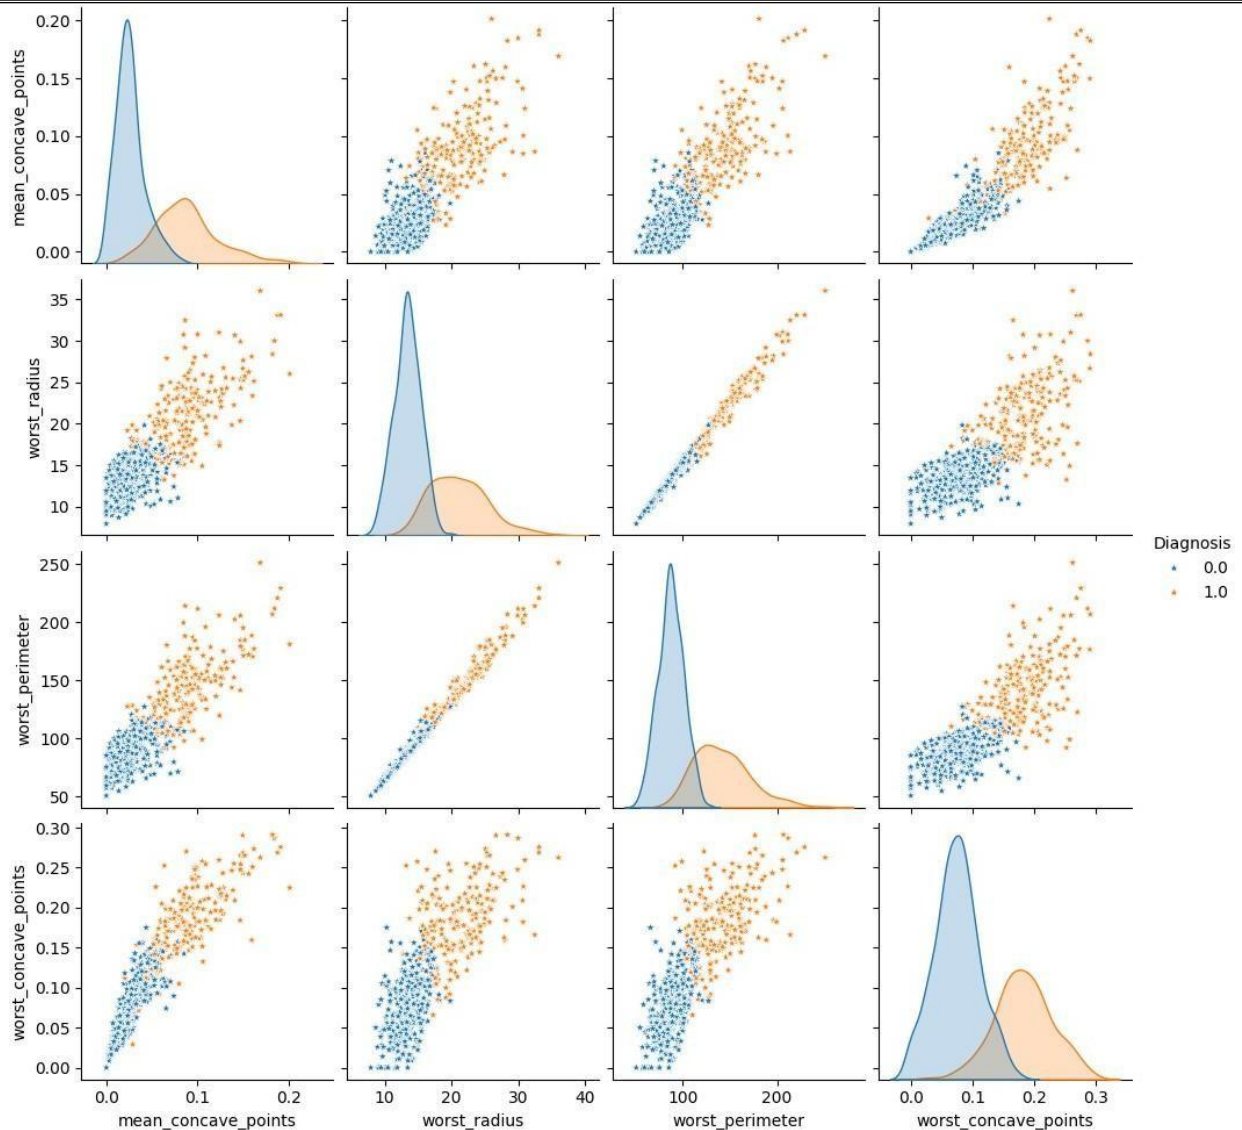
```
# Count occurrences of each unique value in the 'Diagnosis' column  
diagnosis_counts = data['Diagnosis'].value_counts()  
  
# Define colors for each slice  
colors = ['skyblue', 'salmon']  
  
# Plotting the pie chart with custom colors  
plt.figure(figsize=(8, 8))  
plt.pie(diagnosis_counts, labels=diagnosis_counts.index,  
autopct='%1.1f%%', startangle=140, colors=colors)  
plt.title('Distribution of Diagnosis')  
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a
```

```
circle.  
plt.show()
```



1. PAIRPLOT FOR HIGHLY CORRELATED FEATURES

```
threshold = 0.75  
filtre = np.abs(correlation_matrix["Diagnosis"] > threshold)  
corr_features = correlation_matrix.columns[filtre].tolist()  
sns.pairplot(data[corr_features], diag_kind = "kde" , markers = "*",  
hue="Diagnosis")  
plt.show()
```



```
# Remove 'ID' column
data.drop('ID', axis=1, inplace=True)

# Convert 'Diagnosis' column to float64
data['Diagnosis'] = data['Diagnosis'].astype(float)

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
0   Diagnosis                             569 non-null    float64
1   mean_radius                           569 non-null    float64
2   mean_texture                           569 non-null    float64
```



```

3   mean_perimeter      569 non-null    float64
4   mean_area           569 non-null    float64
5   mean_smoothness     569 non-null    float64
6   mean_compactness    569 non-null    float64
7   mean_concavity      569 non-null    float64
8   mean_concave_points 569 non-null    float64
9   mean_symmetry       569 non-null    float64
10  mean_fractal_dimension 569 non-null    float64
11  se_radius           569 non-null    float64
12  se_texture          569 non-null    float64
13  se_perimeter        569 non-null    float64
14  se_area             569 non-null    float64
15  se_smoothness       569 non-null    float64
16  se_compactness      569 non-null    float64
17  se_concavity        569 non-null    float64
18  se_concave_points   569 non-null    float64
19  se_symmetry         569 non-null    float64
20  se_fractal_dimension 569 non-null    float64
21  worst_radius        569 non-null    float64
22  worst_texture       569 non-null    float64
23  worst_perimeter     569 non-null    float64
24  worst_area          569 non-null    float64
25  worst_smoothness    569 non-null    float64
26  worst_compactness   569 non-null    float64
27  worst_concavity     569 non-null    float64
28  worst_concave_points 569 non-null    float64
29  worst_symmetry      569 non-null    float64
30  worst_fractal_dimension 569 non-null    float64
dtypes: float64(31)
memory usage: 137.9 KB

```

SPLITTING DATA INTO TRAIN , TEST AND VALIDATION SETS

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc
from sklearn.ensemble import BaggingClassifier,
RandomForestClassifier, AdaBoostClassifier
import matplotlib.pyplot as plt

# Step 5: Split the data into training, testing, and validation sets
X = data.drop('Diagnosis', axis=1)
y = data['Diagnosis']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train,
test_size=0.25, random_state=42) # 60% train, 20% validation, 20%
test

```

```

from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 10, criterion =
'entropy', random_state = 0)
classifier.fit(X_train, y_train)

RandomForestClassifier(criterion='entropy', n_estimators=10,
random_state=0)

```

TRAINING AND TESTING MODEL

```

#Train the model
models = {
    "Bagging": BaggingClassifier(),
    "Random Forest": RandomForestClassifier(),
    "AdaBoost": AdaBoostClassifier()
}

for name, model in models.items():
    model.fit(X_train, y_train)

#Test the model
results = {}
for name, model in models.items():
    y_pred = model.predict(X_test)
    results[name] = y_pred

c:\Users\HP\AppData\Local\Programs\Python\Python311\Lib\site-packages\
sklearn\ensemble\_weight_boosting.py:519: FutureWarning: The SAMME.R
algorithm (the default) is deprecated and will be removed in 1.6. Use
the SAMME algorithm to circumvent this warning.
    warnings.warn(

# Step 8: Measure the performance of the trained model
# (Assuming binary classification)
# (Assuming binary classification)
def calculate_roc(y_true, y_pred):
    fpr, tpr, thresholds = roc_curve(y_true, y_pred)
    roc_auc = auc(fpr, tpr)
    return fpr, tpr, roc_auc

plt.figure(figsize=(8, 6))

<Figure size 800x600 with 0 Axes>

# Step 9: Compare the results of each ensemble model using graphs
for name, y_pred in results.items():
    fpr, tpr, roc_auc = calculate_roc(y_test, y_pred)
    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

    accuracy = accuracy_score(y_test, y_pred)

```


```

print(f"{name}: Accuracy = {accuracy:.4f}")

plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()

```

Bagging: Accuracy = 0.9649
 Random Forest: Accuracy = 0.9649
 AdaBoost: Accuracy = 0.9561



```

from sklearn.metrics import confusion_matrix
import seaborn as sns

# Step 9: Compare the results of each ensemble model using graphs
for name, y_pred in results.items():

    accuracy = accuracy_score(y_test, y_pred)
    print(f"{name}: Accuracy = {accuracy:.4f}")

```

```

# Print classification report
print(f"\n{name} Classification Report:\n{classification_report(y_test, y_pred)}")


# Plot confusion matrix
plt.figure()
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.title(f'Confusion Matrix - {name}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```

Bagging: Accuracy = 0.9649

Bagging Classification Report:


	precision	recall	f1-score	support
0.0	0.96	0.99	0.97	71
1.0	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114



Random Forest: Accuracy = 0.9649

Random Forest Classification Report:


	precision	recall	f1-score	support
0.0	0.96	0.99	0.97	71
1.0	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114



AdaBoost: Accuracy = 0.9561

AdaBoost Classification Report:


	precision	recall	f1-score	support
0.0	0.96	0.97	0.97	71
1.0	0.95	0.93	0.94	43
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114



```
# Step 10: Represent the ROC of training and test results in the
graphs
plt.figure(figsize=(8, 6))
for name, model in models.items():
    y_train_pred = model.predict(X_train)
    fpr_train, tpr_train, roc_auc_train = calculate_roc(y_train,
y_train_pred)
    plt.plot(fpr_train, tpr_train, label=f'{name} Train (AUC =
{roc_auc_train:.2f})', linestyle='--')

    fpr_test, tpr_test, roc_auc_test = calculate_roc(y_test,
results[name])
    plt.plot(fpr_test, tpr_test, label=f'{name} Test (AUC =
{roc_auc_test:.2f})')

plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic - Train vs Test')
plt.legend(loc='lower right')
plt.show()
```



INFERENCE :

- Bagging and Random Forest models achieved identical accuracy scores of 0.9649, while AdaBoost had a slightly lower accuracy of 0.9561.
- Precision, recall, and F1-score metrics indicate high performance across all models, suggesting strong predictive capability.
- Detailed classification reports provide insights into the performance of each model for both classes (0 and 1), demonstrating their ability to correctly classify instances.
- Overall, Bagging and Random Forest models slightly outperform AdaBoost in diagnosing breast cancer.
- The models exhibit high accuracy and robustness, as evidenced by the AUC values.
- Visualizations of ROC curves can offer further insights into the comparative performance of the models.

LEARNING OUTCOMES :

- Ensemble Learning: Understanding and applying ensemble techniques for classification tasks.
- Data Preprocessing: Handling missing data, encoding categorical variables, and scaling features for model training.
- Model Evaluation: Using classification metrics to assess model performance.
- Feature Engineering: Implementing techniques to enhance model predictive capability.
- Model Evaluation: Familiarity with classification metrics like accuracy, precision, recall, and F1-score for assessing model performance.

GITHUB LINK :

<https://github.com/KavinSiva13/ML-A4/tree/main/Assignment-8>